

# Advanced C and System Programming

Anandkumar



# Threads

- Threads are mechanisms to do more than one job at a time.
- Threads are finer-grained units of execution.
- Threads, unlike processes, share the same address space and other resources.
- POSIX standard thread API is not included in standard C library, they are in *libpthread.so*.
- In Linux, threads are handled by LWPs.



# Threads

A Thread is an independent stream of instructions that can be schedule to run as such by the OS.

Think of a thread as a “procedure” that runs independently from its main program.

Multi-threaded programs are where several procedures are able to be scheduled to run simultaneously and/or independently by the OS.


A Thread exists within a process and uses the process resources.

# Threads (cont)

Threads only duplicate the essential resources it needs to be independently schedulable.

A thread will die if the parent process dies.

A thread is “lightweight” because most of the overhead has already been accomplished through the creation of the process.



# POSIX Threads (PThreads)

For UNIX systems, implementations of threads that adhere to the IEEE POSIX 1003.1c standard are Pthreads.

Pthreads are C language programming types defined in the `pthread.h` header/include file.

# Why Use Pthreads

The primary motivation behind Pthreads is improving program performance.

Can be created with much less OS overhead.

Needs fewer system resources to run.

View comparison of forking processes to using a `pthread_create` subroutine. Timings reflect 50,000 processes/thread creations.

# Threads vs Forks

PLATFORM	fork()			pthread_create()		
	REAL	USER	SYSTEM	REAL	USER	SYSTEM
AMD 2.4 GHz Opteron (8cpus/node)	41.07	60.08	9.01	0.66	0.19	0.43
IBM 1.9 GHz POWER5 p5-575 (8cpus/node)	64.24	30.78	27.68	1.75	0.69	1.1
IBM 1.5 GHz POWER4 (8cpus/node)	104.05	48.64	47.21	2.01	1	1.52
INTEL 2.4 GHz Xeon (2 cpus/node)	54.95	1.54	20.78	1.64	0.67	0.9
INTEL 1.4 GHz Itanium2 (4 cpus/node)	54.54	1.07	22.22	2.03	1.26	0.67

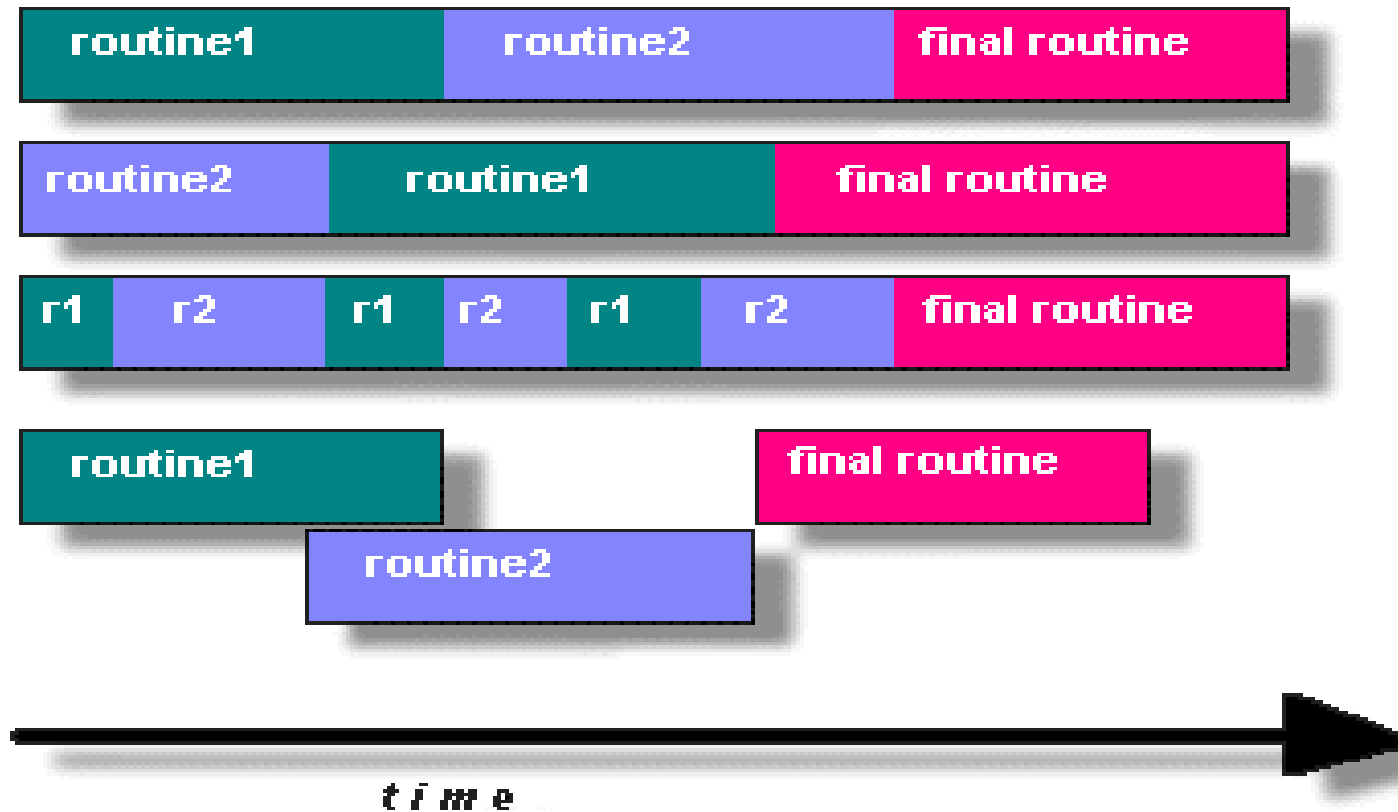


# Designing Pthreads Programs

Pthreads are best used with programs that can be organized into discrete, independent tasks which can execute concurrently.

Example: routine 1 and routine 2 can be interchanged, interleaved and/or overlapped in real time.

# Candidates for Pthreads



# Designing Pthreads (cont)

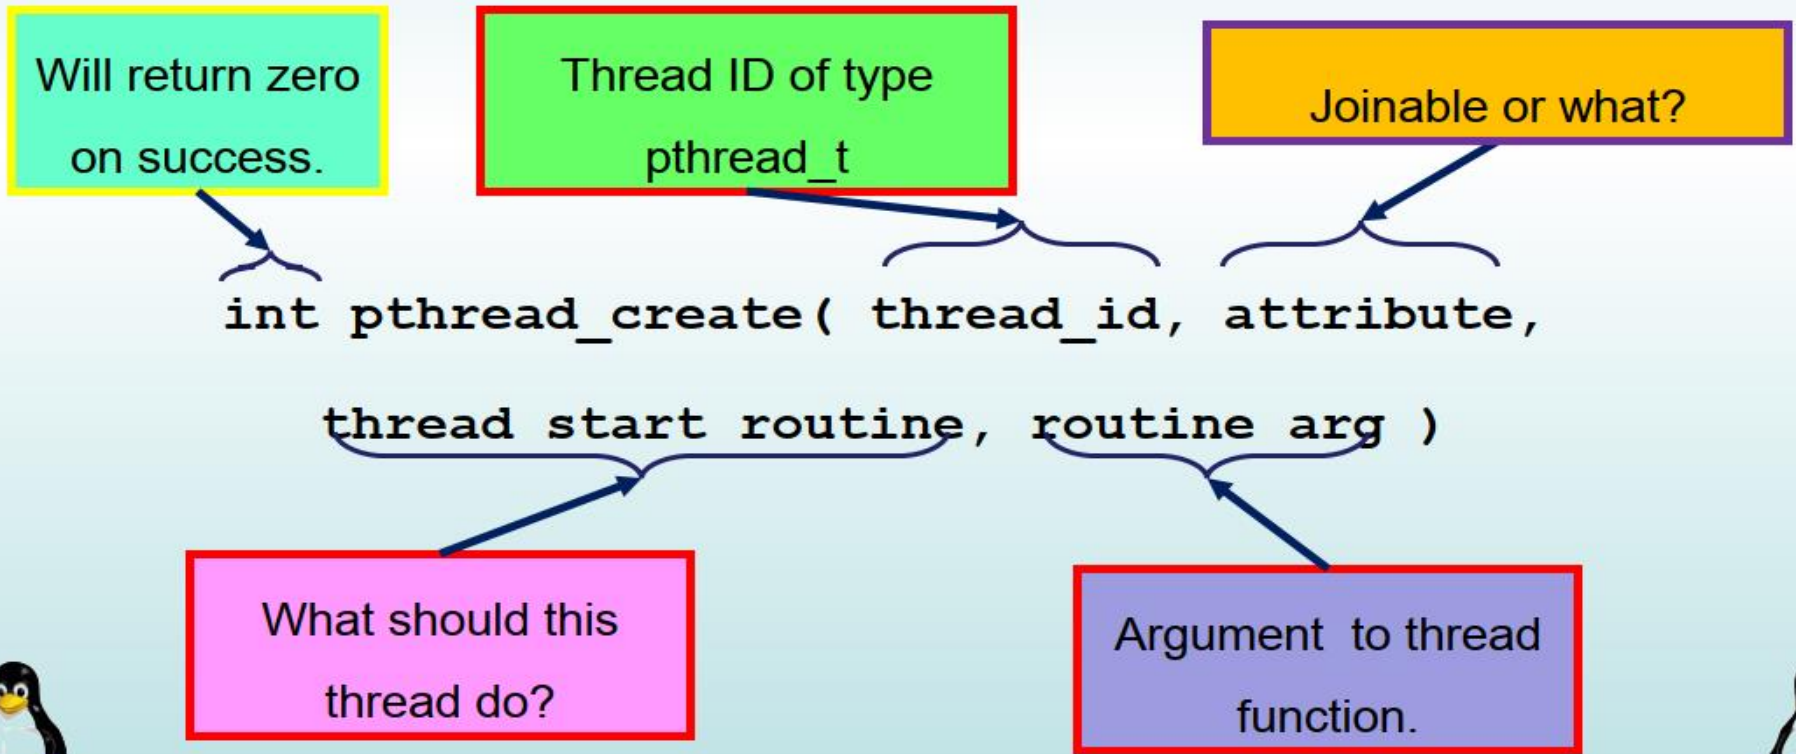
Common models for threaded programs:

Manager/Worker: manager assigns work to other threads, the workers. Manager handles input and hands out the work to the other tasks.

Pipeline: task is broken into a series of suboperations, each handled in series but concurrently, by a different thread.

# Creating threads

- Like processes, each thread has its own Thread-ID of type *pthread\_t*.
- You can create a thread by calling the *pthread\_create* function.



# Creating threads

- *pthread\_create* returns immediately and the specified thread will do its job separately.
- If one of the threads in a program, call *exec* the whole process image will be replaced.
- The argument passed to the thread routine is a *void \**.
- You can pass more data in a structure of type *void \**.



# Pthread Management - Creating Threads

The `main()` method comprises a single, default thread.

`pthread_create()` creates a new thread and makes it executable.

The maximum number of threads that may be created by a process is implementation dependent.

Once created, threads are peers, and may create other threads.

# Joining threads

- You can wait for a thread to finish its job using *pthread\_join*.
- *pthread\_join* is something similar to *wait* function in processes.
- Using *pthread\_join*, you can also take the return value of a thread.
- A thread, can not call *pthread\_join* to wait for itself, you can use *pthread\_self* function to get the TID of running thread and deciding what to do.



# Joining threads

- Like processes, you can wait for a thread to finish its job...

```
int pthread_join( pthread_t thread_id, void ** return_value )
```

Will return zero  
on success.

Thread ID which you  
want to wait for.

The return value of  
thread will be put here.





# Pthread Management - Terminating Threads

Several ways to terminate a thread:

- The thread is complete and returns

- The `pthread_exit()` method is called

- The `pthread_cancel()` method is invoked

- The `exit()` method is called

The `pthread_exit()` routine is called after a thread has completed its work and it no longer is required to exist.

# Terminating Threads (cont)

If the main program finishes before the thread(s) do, the other threads will continue to execute if a `pthread_exit()` method exists.

The `pthread_exit()` method does not close files; any files opened inside the thread will remain open, so cleanup must be kept in mind.

# Pthread Example

```
#include <pthread.h>
```

```
#include <stdio.h>
```

```
#define NUM_THREADS 5
```

```
void *PrintHello(void *threadid)
```

```
{
```

```
    int tid; tid = (int)threadid;
```

```
    printf("Hello World! It's me, thread #%d!\n", tid);
```

```
    pthread_exit(NULL);
```

```
}
```

# Pthread Example

```
int main (int argc, char *argv[])
{
    pthread_t threads[NUM_THREADS];
    int rc, t;
    for(t=0; t<NUM_THREADS; t++)
    {
        printf("In main: creating thread %d\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)t);
        if (rc)
        {
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }
    pthread_exit(NULL);
}
```

# Pthread Example - Output

In main: creating thread 0

In main: creating thread 1

Hello World! It's me, thread #0!

In main: creating thread 2

Hello World! It's me, thread #1!

Hello World! It's me, thread #2!

In main: creating thread 3

In main: creating thread 4

Hello World! It's me, thread #3!

# Thread attributes

- Second parameter in *pthread\_create* is the thread attribute.
- Most useful attribute of a thread is *joinability*.
- If a thread is *joinable*, it is not automatically cleaned up.
- To clean up a *joinable* like a child process, you should call *pthread\_join* .
- A *detached* thread, is automatically cleaned up.
- A joinable thread may be turned into a detached one, but can not be made joinable again.
- Using *pthread\_detach* you can turn a joinable thread into detached.



# Thread attributes

- If you do not clean up the joinable thread, it will become something like zombie.
- To assign an attribute to a thread, you should:
  - Create a *pthread\_attr\_t* object.
  - Call *pthread\_attr\_init* to initialize the attribute object.
  - Modify the attributes.
  - Pass a pointer to *pthread\_create*.
  - Call *pthread\_attr\_destroy* to release the attribute object.



# Thread cancelation

- A thread might be terminated by finishing its job or calling *pthread\_exit* or by a request from another thread.
- The latter case is called “Thread Cancelation”.
- You can cancel a thread using *pthread\_cancel*.
- If the canceled thread is not detached, you should join it after cancelation, otherwise it will become zombie.
- You can disable cancelation of a thread using *pthread\_setcancelstate()*.





# Thread cancelation

- There are two cancel state:
- **PTHREAD\_CANCEL\_ASYNCHRONOUS**: Asynchronously cancelable (cancel at any point of execution)
- **PTHREAD\_CANCEL\_DEFERRED**: Synchronously cancelable (thread checks for cancellation requests)
- There are two cancelation types:
- **PTHREAD\_CANCEL\_DISABLE** and **PTHREAD\_CANCEL\_ENABLE**.
- It's a good idea to set the state to *Uncancelable* when entering critical section...





# Thread-Local Storage

---

- **Thread-local storage (TLS)** allows each thread to have its own copy of data
- Useful when you do not have control over the thread creation process (i.e., when using a thread pool)
- Different from local variables
  - Local variables visible only during single function invocation
  - TLS visible across function invocations
- Similar to `static` data
  - TLS is unique to each thread

# Process

Code

Global Variables

Process Heap

Process Resources

*Open Files*

*Heaps...*

Environment Block

Thread 1

Thread N

Thread Local Storage

...

Thread Local Storage

Stack

Stack

# Critical Section

- The ultimate cause of most bugs involving threads is that they are accessing the same data at the same time.
- The section of code which is responsible to access the shared data, is called *Critical Section* .
- A critical section is part of code that should be executed completely or not at all (a thread should not be interrupted when it is in this section)
- If you do not protect the *Critical Section*, your program might crash because of *Race Condition*.





# Race Condition

- Race Condition is a condition in which threads are racing each other to change the same data structure.
- Because there is no way to know when the system scheduler will interrupt one thread and execute the other one, the buggy program may crash once and finish regularly next time.
- To eliminate race conditions, you need a way to make operations *atomic* (uninterruptible).



# Synchronization Primitives

## Counting Semaphores

Permit a limited number of threads to execute a section of the code

## Binary Semaphores - Mutexes

Permit only one thread to execute a section of the code

## Condition Variables

Communicate information about the state of shared data

# POSIX Semaphores

## Named Semaphores

Provides synchronization between unrelated process and related process as well as between threads

Kernel persistence

→ System-wide and limited in number

Uses `sem_open`

## Unnamed Semaphores

Provides synchronization between threads and between related processes

Thread-shared or process-shared

Uses `sem_init`

# POSIX Semaphores

## Data type

Semaphore is a variable of type `sem_t`

Include `<semaphore.h>`

## Atomic Operations

```
int sem_init(sem_t *sem, int pshared,  
             unsigned value);
```

```
int sem_destroy(sem_t *sem);
```

```
int sem_post(sem_t *sem);
```

```
int sem_trywait(sem_t *sem);
```

```
int sem_wait(sem_t *sem);
```



# Unnamed Semaphores

```
#include <semaphore.h>
```

```
int sem_init(sem_t *sem, int pshared,  
             unsigned value);
```

Initialize an unnamed semaphore

Returns

0 on success

-1 on failure, sets **errno**

Parameters

**sem:**

Target semaphore

**pshared:**

0: only threads of the creating process can use the semaphore

Non-0: other processes can use the semaphore

**value:**

Initial value of the semaphore

You cannot make a copy of a semaphore variable!!!

# Sharing Semaphores

Sharing semaphores between threads within a process is easy, use **pshared==0**

A non-zero **pshared** allows any process that can access the semaphore to use it

Places the semaphore in the global (OS) environment

Forking a process creates copies of any semaphore it has

Note: unnamed semaphores are not shared across unrelated processes

# `sem_init` can fail

On failure

`sem_init` returns -1 and sets `errno`

<code>errno</code>	cause
<code>EINVAL</code>	<code>Value &gt; sem_value_max</code>
<code>ENOSPC</code>	Resources exhausted
<code>EPERM</code>	Insufficient privileges

```
sem_t semA;
```

```
if (sem_init(&semA, 0, 1) == -1)
    perror("Failed to initialize semaphore
semA");
```

# Semaphore Operations

```
#include <semaphore.h>
```

```
int sem_destroy(sem_t *sem) ;
```

Destroy an semaphore

Returns

0 on success

-1 on failure, sets **errno**

Parameters

**sem:**

Target semaphore

Notes

Can destroy a **sem\_t** only once

Destroying a destroyed semaphore gives undefined results

Destroying a semaphore on which a thread is blocked gives undefined results

# Semaphore Operations

```
#include <semaphore.h>
```

```
int sem_post(sem_t *sem) ;
```

Unlock a semaphore - same as signal

Returns

0 on success

-1 on failure, sets `errno` (`== EINVAL` if semaphore doesn't exist)

Parameters

`sem`:

Target semaphore

`sem > 0`: no threads were blocked on this semaphore, the semaphore value is incremented

`sem == 0`: one blocked thread will be allowed to run

# Semaphore Operations

```
#include <semaphore.h>
```

```
int sem_wait(sem_t *sem);
```

Lock a semaphore

Blocks if semaphore value is zero

Returns

0 on success

-1 on failure, sets **errno** (**== EINTR** if interrupted by a signal)

Parameters

**sem:**

Target semaphore

sem > 0: thread acquires lock

sem == 0: thread blocks

# Semaphore Operations

```
#include <semaphore.h>
```

```
int sem_trywait(sem_t *sem);
```

Test a semaphore's current condition

Does not block

Returns

0 on success

-1 on failure, sets `errno` (`== AGAIN` if semaphore already locked)

Parameters

**sem:**

Target semaphore

`sem > 0`: thread acquires lock

`sem == 0`: thread returns

# Pthread Mutex

## States

Locked

Some thread  
holds  
the  
mutex

Unlocked

No thread  
holds  
the  
mutex

When several threads  
compete

One wins

The rest block

Queue of  
blocked  
threads



# Mutex Variables

A typical sequence in the use of a mutex

1. Create and initialize **mutex**
2. Several threads attempt to lock **mutex**
3. Only one succeeds and now owns **mutex**
4. The owner performs some set of actions
5. The owner unlocks **mutex**
6. Another thread acquires **mutex** and repeats the process
7. Finally **mutex** is destroyed

# Creating a mutex

```
#include <pthread.h>
```

```
int pthread_mutex_init(pthread_mutex_t *mutex,  
    const pthread_mutexattr_t *attr);
```

Initialize a pthread mutex: the mutex is initially unlocked

Returns

0 on success

Error number on failure

**EAGAIN**: The system lacked the necessary resources; **ENOMEM**: Insufficient memory ; **EPERM**: Caller does not have privileges; **EBUSY**: An attempt to re-initialise a mutex; **EINVAL**: The value specified by attr is invalid

Parameters

**mutex**: Target mutex

**attr**:

NULL: the default mutex attributes are used

Non-NULL: initializes with specified attributes

# Creating a mutex

## Default attributes

Use `PTHREAD_MUTEX_INITIALIZER`

Statically allocated

Equivalent to dynamic initialization by a call to `pthread_mutex_init()` with parameter `attr` specified as NULL

No error checks are performed

# Destroying a mutex

```
#include <pthread.h>

int
pthread_mutex_destroy(pthread_mutex_t
*mutex);
```

Destroy a pthread mutex

Returns

0 on success

Error number on failure

**EBUSY**: An attempt to re-initialise a mutex; **EINVAL**: The value specified by attr is invalid

Parameters

**mutex**: Target mutex

# Locking/unlocking a mutex

```
#include <pthread.h>

int pthread_mutex_lock(pthread_mutex_t
    *mutex) ;

int
    pthread_mutex_trylock(pthread_mutex_t
    *mutex) ;

int pthread_mutex_unlock(pthread_mutex_t
    *mutex) ;
```

## Returns

0 on success

Error number on failure

**EBUSY**: already locked; **EINVAL**: Not an initialised mutex;

**EDEADLK**: The current thread already owns the mutex;

**EPERM**: The current thread does not own the mutex

# Simple Example

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

static pthread_mutex_t my_lock =
    PTHREAD_MUTEX_INITIALIZER;

void *mythread(void *ptr) {
    long int i,j;
    while (1) {
        pthread_mutex_lock (&my_lock);

        for (i=0; i<10; i++) {
            printf ("Thread %d\n", (int) ptr);
            for (j=0; j<50000000; j++);
        }
    }
}
```

```
int main (int argc, char *argv[]) {
    pthread_t thread[2];

    pthread_create(&thread[0], NULL,
        mythread, (void *)0);

    pthread_create(&thread[1], NULL,
        mythread, (void *)1);

    getchar();
}
```

# Condition Variables

Used to communicate information  
about the state of shared data

Execution of code depends on the state  
of

- A data structure or

- Another running thread

Allows threads to synchronize based  
upon the actual value of data

Without condition variables

Threads continually poll to check if the  
condition is met

# Condition Variables

## Signaling, not mutual exclusion

A mutex is needed to synchronize access to the shared data

Each condition variable is associated with a single mutex

Wait atomically unlocks the mutex and blocks the thread

Signal awakens a blocked thread



# Creating a Condition Variable

Similar to pthread mutexes

```
int pthread_cond_init(pthread_cond_t  
    *cond, const pthread_condattr_t  
    *attr);
```

```
int pthread_cond_destroy(pthread_cond_t  
    *cond);
```

```
pthread_cond_t cond =  
    PTHREAD_COND_INITIALIZER;
```

# Using a Condition Variable

## Waiting

Block on a condition variable.

Called with **mutex** locked by the calling thread

Atomically release **mutex** and cause the calling thread to block on the condition variable

On return, **mutex** is locked again

```
int pthread_cond_wait(pthread_cond_t *cond,  
    pthread_mutex_t *mutex);
```

```
int pthread_cond_timedwait(pthread_cond_t *cond,  
    pthread_mutex_t *mutex, const struct timespec  
    *abstime);
```

# Using a Condition Variable

## Signaling

```
int pthread_cond_signal(pthread_cond_t *cond);
```

unblocks at least one of the blocked threads

```
int pthread_cond_broadcast(pthread_cond_t *cond);
```

unblocks all of the blocked threads

## Signals are not saved

Must have a thread waiting for the signal or it will be lost

# Spinlock

Spin locks are a low-level synchronization mechanism suitable primarily for use on shared memory multiprocessors. When the calling thread requests a spin lock that is already held by another thread, the second thread spins in a loop to test if the lock has become available. When the lock is obtained, it should be held only for a short time, as the spinning wastes processor cycles. Callers should unlock spin locks before calling sleep operations to enable other threads to obtain the lock.

# Spinlock

## pthread\_spin\_init() Syntax

```
int pthread_spin_init(pthread_spinlock_t *lock, int pshared);
```

```
#include <pthread.h>

pthread_spinlock_t lock;
int pshared;
int ret;

/* initialize a spin lock */
ret = pthread_spin_init(&lock, pshared);
```

The *pshared* attribute has one of the following values:

PTHREAD\_PROCESS\_SHARED

**Description:** Permits a spin lock to be operated on by any thread that has access to the memory where the spin lock is allocated. Operation on the lock is permitted even if the lock is allocated in memory that is shared by multiple processes.

PTHREAD\_PROCESS\_PRIVATE

**Description:** Permits a spin lock to be operated upon only by threads created within the same process as the thread that initialized the spin lock. If threads of differing processes attempt to operate on such a spin lock, the behavior is undefined. The default value of the process-shared attribute is PTHREAD\_PROCESS\_PRIVATE.

# Spinlock

## pthread\_spin\_lock() Syntax

```
int pthread_spin_lock(pthread_spinlock_t *lock);
```

```
#include <pthread.h>
```

```
pthread_spinlock_t lock;  
int ret;
```

```
ret = pthread_spin_lock(&lock); /* lock the spinlock */
```

## pthread\_spin\_unlock() Syntax

```
int pthread_spin_unlock(pthread_spinlock_t *lock);
```

```
#include <pthread.h>
```

```
pthread_spinlock_t lock;  
int ret;
```

```
ret = pthread_spin_unlock(&lock); /* spinlock is unlocked */
```

## pthread\_spin\_destroy() Syntax

```
int pthread_spin_destroy(pthread_spinlock_t *lock);
```

```
#include <pthread.h>
```

```
pthread_spinlock_t lock;  
int ret;
```

```
ret = pthread_spin_destroy(&lock); /* spinlock is destroyed */
```

# RW lock

Operation	Related Function Description
Initialize a read-write lock	<code>pthread_rwlock_init</code>
Read lock on read-write lock	<code>pthread_rwlock_rdlock</code>
Read lock with a nonblocking read-write lock	<code>pthread_rwlock_tryrdlock</code>
Write lock on read-write lock	<code>pthread_rwlock_wrlock</code>
Write lock with a nonblocking read-write lock	<code>pthread_rwlock_trywrlock</code>
Unlock a read-write lock	<code>pthread_rwlock_unlock</code>
Destroy a read-write lock	<code>pthread_rwlock_destroy</code>

# RW lock

## `pthread_rwlock_init` Syntax

```
#include <pthread.h>

int pthread_rwlock_init(pthread_rwlock_t *restrict rwlock,
                        const pthread_rwlockattr_t *restrict attr);

pthread_rwlock_t rwlock = PTHREAD_RWLOCK_INITIALIZER;
```

## `pthread_rwlock_wrlock` Syntax

```
#include <pthread.h>

int pthread_rwlock_wrlock(pthread_rwlock_t *rwlock );
```

## `pthread_rwlock_destroy` Syntax

```
#include <pthread.h>

int pthread_rwlock_destroy(pthread_rwlock_t **rwlock);
```

## `pthread_rwlock_rdlock` Syntax

```
#include <pthread.h>

int pthread_rwlock_rdlock(pthread_rwlock_t *rwlock );
```

## `pthread_rwlock_unlock` Syntax

```
#include <pthread.h>

int pthread_rwlock_unlock (pthread_rwlock_t *rwlock);
```



# Thread Safe

```
...  
...  
...  
  
char arr[10];  
int index=0;  
  
int func(char c)  
{  
    int i=0;  
    if(index >= sizeof(arr))  
    {  
        printf("\n No storage\n");  
        return -1;  
    }  
    arr[index] = c;  
    index++;  
    return index;  
}  
  
...  
...  
...
```

# Thread Safe

```
char arr[10];
int index=0;

int func(char c)
{
    int i=0;
    if(index >= sizeof(arr))
    {
        printf("\n No storage\n");
        return -1;
    }

    /* ...
       Lock a mutex here
       ...
    */

    arr[index] = c;
    index++;

    /* ...
       unlock the mutex here
       ...
    */

    return index;
}
```

# Producers Consumers Systems

One system produce items that will be used by other system

## Examples

shared printer, the printer here acts the consumer, and the computers that produce the documents to be printed are the consumers.

Sensors network, where the sensors here the producers, and the base stations (sink) are the producers.

# Producer Consumer Problem

The producer-consumer problem illustrates the need for synchronization in systems where many processes share a resource. In the problem, two processes share a fixed-size buffer. One process produces information and puts it in the buffer, while the other process consumes information from the buffer. These processes do not take turns accessing the buffer, they both work concurrently.

It is also called bounded buffer problem

# Producer

```
While(Items_number ==buffer size)  
    ; //waiting since the buffer is full
```

```
Buffer[i]=next_produced_item;  
i=(i+1)%Buffer_size;  
Items_number++;
```

# Consumer

```
while (Items_number == 0)
```

```
    ; // do nothing since the buffer is  
    empty
```

```
        Consumed _item= buffer[j];
```

```
        j = (j + 1) % Buffer_size;
```

```
        Items_number--;
```